

# Introduction to the R programming language



Dan Sewell  
Department of Biostatistics, University of Iowa

- 1 Basics of the Programming Language
  - Introduction
  - Data Structures
  - Programming Structures
  - Reading/Writing Data
- 2 Linear Algebra
- 3 Data visualization
  - Univariate
  - Multivariate
  - ggplot2
  - 3D plotting
- 4 (Legal) Performance Enhancers
- 5 Extended Examples
  - Blowdown data
  - UK monthly deaths from lung diseases

# Introduction to R

# What is R?

From the R Core Team:

“R is a system for statistical computation and graphics.”

- high level programming language
- run-time environment with graphics
- debugger
- data management tools
- sophisticated analytic tools
- thousands of packages are available
- open source- you can contribute!

# Getting Started

## Download R from Comprehensive R Archive Network (CRAN)

[www.r-project.com](http://www.r-project.com)

### The R Project for Statistical Computing

#### Getting Started

R is a free software environment for statistical computing and graphics. It compiles and runs on a wide variety of UNIX platforms, Windows and MacOS. To [download R](#), please choose your preferred [CRAN mirror](#).

If you have questions about R like how to download and install the software, or what the license terms are, please read our [answers to frequently asked questions](#) before you send an email.

#### News

- The [R Logo](#) is available for download in high-resolution PNG or SVG formats.
- [useR! 2016](#), will take place at Stanford University, CA, USA, June 27 - June 30, 2016.
- [The R Journal Volume 7/2](#) is available.
- [R version 3.2.3 \(Wooden Christmas-Tree\)](#) has been released on 2015-12-10.
- [R version 3.1.3 \(Smooth Sidewalk\)](#) has been released on 2015-03-09.



[\[Home\]](#)

**Download**

[CRAN](#)

**R Project**

[About R](#)

[Logo](#)

[Contributors](#)

[What's New?](#)

[Mailing Lists](#)

[Bug Tracking](#)

[Conferences](#)

[Search](#)

# RStudio

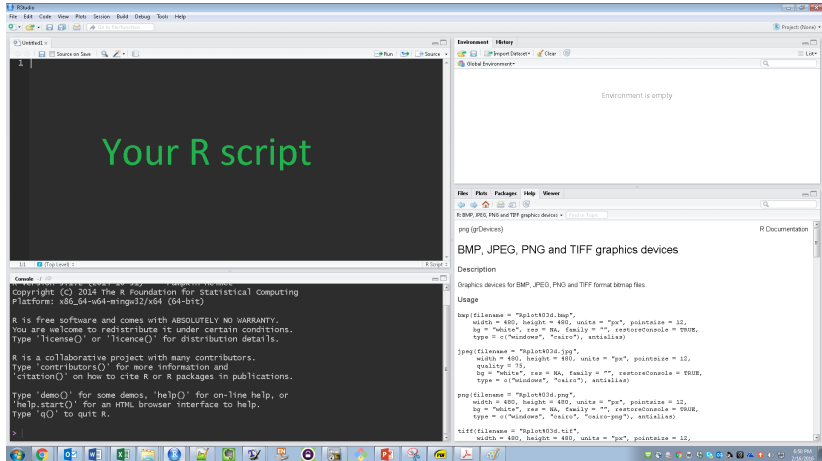
RStudio IDE (Integrated Development Environment)

[www.rstudio.com](http://www.rstudio.com)

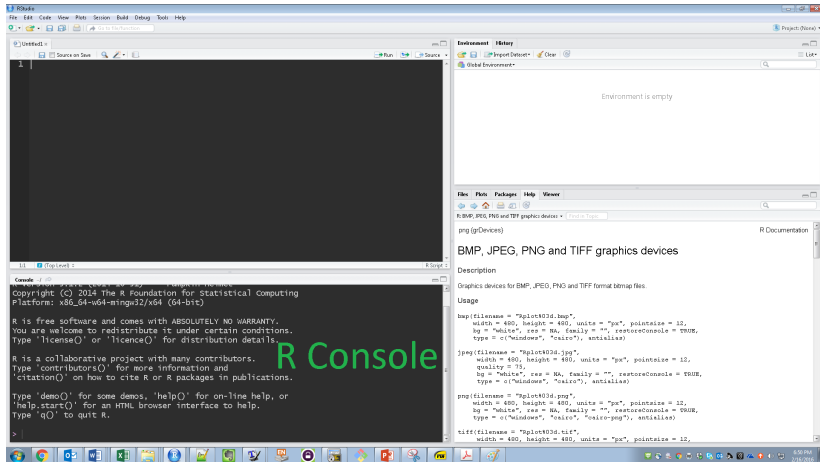
- Nicier GUI
- Lots of features
- Integrated version control
- integration with Sweave/knitr/markdown/Shiny



# RStudio

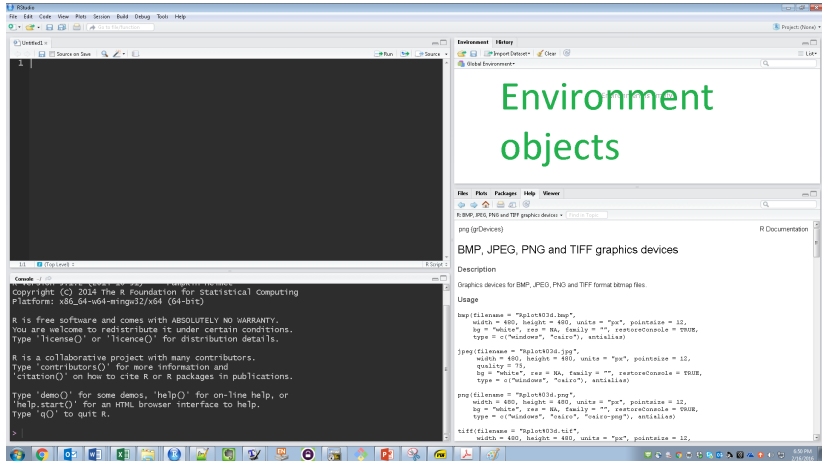


# RStudio

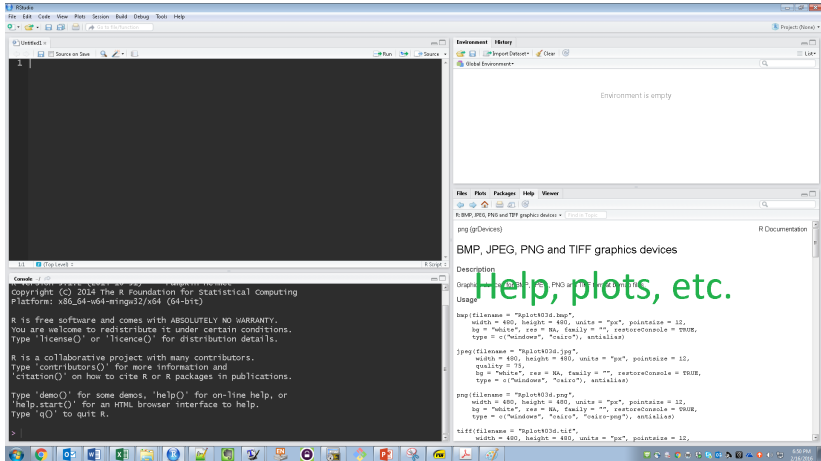




# RStudio



# RStudio



# Scripts

R has a run-time environment, but we can still save code in scripts  
(.R extension), or source them:

```
> setwd("~/testDirectory/")  
> #use getwd() to see working directory  
> source("testScript.R")
```

# Basics

- ; or a newline separate commands

# Basics

- ; or a newline separate commands
- R code is *case sensitive*

# Basics

- ; or a newline separate commands
- R code is *case sensitive*
- # comments all characters until the next newline

# Basics

- ; or a newline separate commands
- R code is *case sensitive*
- # comments all characters until the next newline
- To get help, use ?keyword, or ??searchword, or F1, or TAB

# Workspaces

The R workspace is the collection of all R objects. To view. . .

```
> ls()
```

```
[1] "DataSet" "X"          "Y"
```

To save an object:

```
> save(X, file="test.RData")
```

To save your entire workspace:

```
> save.image("test.RData")
```

To load an R object or R workspace:

```
> load("test.RData")
```

Where are you saving to?

```
> getwd()
```

```
> setwd("~/newdir/")
```



# Packages

The biggest strength of R: PACKAGES

First you must install them (only run once):

```
> #install.packages("lme4")
```

Each R session, load the necessary packages:

```
> library("lme4")
```

```
> #require("lme4")
```

# Packages

The biggest difference between `library` and `require` is which throws an error:

```
> library("asdf");print("script is killed")
Error in library("asdf") : there is no pack-
age called 'asdf'

> require("asdf");print("code is still running!")

[1] "code is still running!"
```

# R script

The accompanying R script for this seminar may be found at  
<http://myweb.uiowa.edu/dksewell/RSeminarScript.R>

# Data Structures

# Classes

- integer

```
> 5L
[1] 5
> as.integer(5.0)
[1] 5
```

# Classes

- integer

```
> 5L
```

```
[1] 5
```

```
> as.integer(5.0)
```

```
[1] 5
```

- numeric

```
> 5.25
```

```
[1] 5.25
```

```
> pi
```

```
[1] 3.141593
```

```
> as.numeric(1L)
```

```
[1] 1
```

# Classes

- character

```
> "hello world"
[1] "hello world"
> as.character(1.0)
[1] "1"
```

# Classes

- character

```
> "hello world"
```

```
[1] "hello world"
```

```
> as.character(1.0)
```

```
[1] "1"
```

- logical/boolean

```
> TRUE
```

```
[1] TRUE
```

```
> 1 > 2
```

```
[1] FALSE
```

```
> as.logical(1L)
```

```
[1] TRUE
```



# Classes

- Missing value

```
> NA
```

```
[1] NA
```

```
> is.na(NA)
```

```
[1] TRUE
```

# Classes

- Missing value

```
> NA
```

```
[1] NA
```

```
> is.na(NA)
```

```
[1] TRUE
```

- NULL

```
> NULL
```

```
NULL
```

```
> is.null(NULL)
```

```
[1] TRUE
```

# Scalars

One can create a scalar object:

```
> x1= 1/3  
> x2 <- 1/6  
> 1/2 -> x3  
> class(x1)  
  
[1] "numeric"
```

# Scalars

One can create a scalar object:

```
> x1= 1/3  
> x2 <- 1/6  
> 1/2 -> x3  
> class(x1)
```

```
[1] "numeric"
```

The R console can also act as a calculator:

```
> (x1+x2-2*x3)*0.1  
[1] -0.05
```

# Vectors

One can create vectors:

```
> v1=c(1,2,4,6)
> v2=1:4
> ( v3=0:3*5+1 )

[1] 1 6 11 16
```

# Vectors

One can create vectors:

```
> v1=c(1,2,4,6)
> v2=1:4
> ( v3=0:3*5+1 )

[1] 1 6 11 16
```

These may also be named

```
> ( v1 = c(x1=1,x2=2,x3=4,x4=6) )

x1 x2 x3 x4
1 2 4 6

> names(v2) = c("x1","x2","x3","x4"); print(v2)

x1 x2 x3 x4
1 2 3 4

> names(v3) = paste("x",1:4,sep=""); print(v3)

x1 x2 x3 x4
1 6 11 16
```

# Vectors

You can retrieve elements or subsets:

```
> v1[4]
```

```
x4  
6
```

```
> v1["x4"]
```

```
x4  
6
```

```
> v1[3:4]
```

```
x3 x4  
4 6
```

```
> v1[c(1,2,4)]
```

```
x1 x2 x4  
1 2 6
```

```
> v1[-c(1,2,4)]
```

```
x3  
4
```

# Vectors

Vectors also may be of varying classes:

```
> ( v4 = c("Hello","World") )
```

```
[1] "Hello" "World"
```

```
> class(v4)
```

```
[1] "character"
```

```
> ( v5 = rep(TRUE,4) )
```

```
[1] TRUE TRUE TRUE TRUE
```

```
> class(v5)
```

```
[1] "logical"
```



# Factors

One may create nominal variables via factors:

```
> ( fac1= factor(rep(1:3,3),labels=paste("Drug",1:3,sep=""))
[1] Drug1 Drug2 Drug3 Drug1 Drug2 Drug3 Drug1 Drug2 Drug3
Levels: Drug1 Drug2 Drug3
```

One may rename the levels as well:

```
> levels(fac1) = paste("Drug",3:1,sep="")
> fac1
[1] Drug3 Drug2 Drug1 Drug3 Drug2 Drug1 Drug3 Drug2 Drug1
Levels: Drug3 Drug2 Drug1
```

# Matrices

One may create matrices (column major by default):

```
> ( mat1 = matrix(0L,2,3) )
```

```
      [,1] [,2] [,3]
[1,]    0    0    0
[2,]    0    0    0
```

```
> ( mat2 = matrix(1:6,2,3) )
```

```
      [,1] [,2] [,3]
[1,]     1     3     5
[2,]     2     4     6
```

```
> ( mat3 = matrix(1:6,2,3,byrow=TRUE) )
```

```
      [,1] [,2] [,3]
[1,]     1     2     3
[2,]     4     5     6
```

# Matrices

And name the rows and columns:

```
> rownames(mat2)=paste("row",1:2)
> colnames(mat2)=paste("column",1:3)
> mat2
```

	column 1	column 2	column 3
row 1	1	3	5
row 2	2	4	6

# Matrices

Again we can extract elements or subsets:

```
> mat2[2,3]
```

```
[1] 6
```

```
> mat2["row 2", "column 3"]
```

```
[1] 6
```

```
> mat2[6]
```

```
[1] 6
```

```
> mat2[,c(1,3)]
```

	column 1	column 3
row 1	1	5
row 2	2	6

# Matrices

There are lots of types of matrices and packages that deal with them, e.g., “Matrix”

```
> require("Matrix")
> mat4 = matrix(0,1000,50)
> mat5 = Matrix(0,1000,50,sparse=TRUE)
> class(mat4)
```

```
[1] "matrix"
```

```
> object.size(mat4)
```

```
400200 bytes
```

```
> class(mat5)
```

```
[1] "dgCMatrix"
```

```
attr(,"package")
```

```
[1] "Matrix"
```

```
> object.size(mat5)
```

```
1624 bytes
```

# Matrices

Other special packages for sparse matrices:

- Matrix
- glmnet
- SparseM
- slam
- spam
- igraph/network
- irlba (Fast and memory efficient methods for truncated singular and eigenvalue decompositions and principal component analysis of large sparse or dense matrices.)

# Matrices

Special packages for huge general matrices:

- bigmemory (Manage Massive Matrices with Shared Memory and Memory-Mapped Files)
- bigalgebra (BLAS routines for native R matrices and big.matrix objects)
- bigtabulate (Table, Apply, and Split Functionality for Matrix and 'big.matrix' Objects)
- biganalytics (Utilities for 'big.matrix' Objects from Package 'bigmemory')
- bigpca (PCA, Transpose and Multicore Functionality for 'big.matrix' Objects)

# Matrices

Using bigmemory:

```
> require("bigmatrix");require("bigalgebra")
> mat6 = matrix(1.0,5000,5000)
> mat7 = as.big.matrix(mat6)
> object.size(mat6);object.size(mat7)
> #200MB vs. 0.6KB
> system.time(mat6%*%mat6)
> #76.44sec
> system.time(mat7%*%mat7)
> #82.03sec
```



# Arrays/Tensors

One may create multidimensional arrays:

```
> ( arr1 = array(1:12,c(2,3,2)) )
```

```
, , 1
```

	[,1]	[,2]	[,3]
[1,]	1	3	5
[2,]	2	4	6

```
, , 2
```

	[,1]	[,2]	[,3]
[1,]	7	9	11
[2,]	8	10	12

# Arrays/Tensors

And extract in exactly the same way as before:

```
> arr1[2,1,1]
```

```
[1] 2
```

```
> arr1[1,,2]
```

```
[1] 7 9 11
```

# Data frames

Data are most often in data frames:

```
> ( df1 = data.frame(id=1:6,
+                    gender=factor(rep(c("M","F"),3)),
+                    treatment1=factor(rep(LETTERS[1:3],2)),
+                    treatment2=factor(rep(LETTERS[1:2],each=3)),
+                    response=rnorm(6)) )
```

	id	gender	treatment1	treatment2	response
1	1	M	A	A	-0.6264538
2	2	F	B	A	0.1836433
3	3	M	C	A	-0.8356286
4	4	F	A	B	1.5952808
5	5	M	B	B	0.3295078
6	6	F	C	B	-0.8204684

**Note:** The columns may be of varying formats!

# Data frames

And extraction:

```
> df1$treatment1
```

```
[1] A B C A B C
```

```
Levels: A B C
```

```
> df1[, "treatment1"]
```

```
[1] A B C A B C
```

```
Levels: A B C
```

```
> df1[,3]
```

```
[1] A B C A B C
```

```
Levels: A B C
```

# Lists

One may create lists of *anything*:

```
> ls1 = list()
> ls1[[1]] = "Hello World"
> ls1[[2]] = pi
> ls1[[3]] = matrix(1:4,2,2)
> print(ls1)
```

```
[[1]]
[1] "Hello World"
```

```
[[2]]
[1] 3.141593
```

```
[[3]]
      [,1] [,2]
[1,]    1    3
[2,]    2    4
```

# Functions

Example: Dose-response Exponential model

```
> DRfun = function(dose){
+   return( 1-exp(-2.18E-04*dose) )
+ }
> DRfun(3.18E+03)
```

```
[1] 0.5000464
```

```
> DRfun(1+03)
```

```
[1] 0.0008716199
```

# Functions

Example: Dose-response Exponential model

```
> DRfun = function(dose,K){  
+   return( 1-exp(-K*dose) )  
+ }  
> DRfun(1e3,2e-4)  
  
[1] 0.1812692
```

## Special extraction commands

Perhaps the most useful is `which()`

```
> df1[which(df1$response<0),]
```

	id	gender	treatment1	treatment2	response
1	1	M	A	A	-0.6264538
3	3	M	C	A	-0.8356286
6	6	F	C	B	-0.8204684



## Special extraction commands

Use `with()` to run commands using a particular data set:

```
> with(df1, response[which(treatment1=="C")])
```

```
[1] -0.8356286 -0.8204684
```

## Special extraction commands

Use `with()` to run commands using a particular data set:

```
> with(df1, response[which(treatment1=="C")])
```

```
[1] -0.8356286 -0.8204684
```

Use `by()` to perform functions according to a grouping variable:

```
> with(df1, by(response, gender, mean))
```

```
gender: F
```

```
[1] 0.3194852
```

```
-----  
gender: M
```

```
[1] -0.3775249
```

# Programming Structures

# for loops

Keywords: for and in

```
> for(i in 1:4){  
+   print(i^2)  
+ }
```

```
[1] 1  
[1] 4  
[1] 9  
[1] 16
```

# for loops

More generally:

```
> x = c(1,2,4,8,16)
> for(i in x){
+   print(log2(i))
+ }
```

```
[1] 0
```

```
[1] 1
```

```
[1] 2
```

```
[1] 3
```

```
[1] 4
```

## while loops

Keywords: while

```
> count=0  
> while(count<5){  
+   print(count)  
+   count = count + 1  
+ }
```

```
[1] 0
```

```
[1] 1
```

```
[1] 2
```

```
[1] 3
```

```
[1] 4
```

## while loops

```
> set.seed(1)
> x=0
> while(x<0.5){
+   x = rnorm(1)
+   print(x)
+ }
```

```
[1] -0.6264538
[1] 0.1836433
[1] -0.8356286
[1] 1.595281
```

## if/else statements

Keywords: if, else, ifelse

```
> set.seed(1)
> uu = runif(1)
> if(uu < 0.5){
+   print("heads")
+ }else{
+   print("tails")
+ }
```

```
[1] "heads"
```



## if/else statements

Keywords: if, else, ifelse

```
> set.seed(1)
> uu = runif(1)
> result = ifelse(uu < 0.5, "heads", "tails")
> print(result)

[1] "heads"
```

# Reading/Writing Data

# Reading in Data

To read data into R:

The screenshot shows the RStudio interface. The top menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Tools, and Help. The left pane shows the Source editor with a script file named '1.R'. The bottom pane shows the Console with a list of data points. The right pane is divided into two sections: 'Environment' and 'Help'. In the 'Environment' section, the 'Import Dataset...' button is circled in red. Below it, the text 'Environment is empty' is visible. The 'Help' section shows the 'R Documentation' for the 'sleepstudy' dataset, titled 'Reaction times in a sleep deprivation study'. The documentation includes a description of the dataset, its format (a data frame with 100 observations), and details about the study.

Environment

Import Dataset...

Environment is empty

R Documentation

Reaction times in a sleep deprivation study

Description

The average reaction time per day for subjects in a sleep deprivation study. On day 0 the subjects had their normal amount of sleep. Starting that night they were restricted to 3 hours of sleep per night. The observations represent the average reaction time on a series of tests given each day to each subject.

Format

A data frame with 100 observations on the following 3 variables.

Reaction

Average reaction time (ms)

Days

Number of days of sleep deprivation

Subject

Subject number on which the observation was made.

Details

These data are from the study described in Beekley et al. (2003), for the sleep-deprived group and for the first 10 days of the study, up to the recovery period.

# Reading in Data

To read data into R:

The screenshot shows the RStudio interface with the following components:

- Select File to Import dialog:** Shows the file 'sleepstudy.csv' (4.11 KB) selected in the 'Desktop' folder. The file name is 'sleepstudy.csv'.
- R Console:** Displays the first 10 rows of the 'sleepstudy' dataset:
 

```
166 166 264.5120 5 371
167 167 259.2658 6 371
168 168 304.6306 7 371
169 169 350.7807 8 371
170 170 369.4692 9 371
171 171 269.4117 0 372
172 172 273.4740 1 372
173 173 297.5968 2 372
174 174 310.6316 3 372
175 175 287.1726 4 372
```
- R Environment:** Shows the 'sleepstudy' object loaded into the environment.
- R Documentation:** Shows the description of the 'sleepstudy' dataset:
 

**Reaction times in a sleep deprivation study**

**Description**

The average reaction time per day for subjects in a sleep deprivation study. On day 0 the subjects had their normal amount of sleep. Shifting that night they were restricted to 3 hours of sleep per night. The observations represent the average reaction time on a series of tests given each day to each subject.

**Format**

A data frame with 100 observations on the following 3 variables.

**Reaction**

Average reaction time (ms)

**Days**

Number of days of sleep deprivation

**Subject**

Subject number on which the observation was made.

**Details**

These data are from the study described in Beekun et al. (2003), for the sleep-deprived group and for the first 10 days of the study, up to the recovery period.

# Reading in Data

## To read data into R:

Environment is empty

Import Dataset

Name: sleepstudy

Input File: [Path to file]

Encoding: Automatic

Header: Yes

Row names: Automatic

Separator: Comma

Decimal: Period

Quote: Double quote

Comment: None

Settings as factors: ☒

Data Frame

X	Reaction	Days	Subject
1	249.1500	1	1
2	248.7000	1	2
3	250.0000	1	3
4	250.0000	1	4
5	251.4000	1	5
6	251.4000	1	6
7	250.0000	1	7
8	250.0000	1	8
9	251.4000	1	9
10	250.0000	1	10

Console

```

166 166 284.5120 5 371
167 167 259.2658 7 371
168 168 304.6306 7 371
169 169 350.7807 8 371
170 170 369.4692 9 371
171 171 269.4117 0 372
172 172 273.4740 1 372
173 173 297.5968 2 372
174 174 310.6316 3 372
175 175 287.1726 4 372
176 176 329.6076 5 372
177 177 334.4818 6 372
178 178 343.2189 7 372
179 179 369.1417 8 372
180 180 364.1236 9 372
  
```

sleep deprivation study

Subjects in a sleep deprivation study. On day 0 the subjects had their normal amount of sleep. Shifting of sleep per night. The observations represent the average reaction time on a series of tests given

Following 3 variables:

Days: Number of days of sleep deprivation

Subject: Subject number on which the observation was made

Details: These data are from the study described in Beekun et al. (2003), for the sleep-deprived group and for the first 15 days of the study, up to the recovery period.

## Reading in Data

```
> sleepstudy = read.csv("http://myweb.uiowa.edu/dksewell/sleepstudy.csv")
> test = read.table("http://myweb.uiowa.edu/dksewell/sleepstudy.csv",
+                  sep=";", header=TRUE)
> all.equal(sleepstudy, test)
```

# Reading in Data

- `require("foreign")` can import Minitab, SAS, Stata, SPSS, etc.
- `require("xlsx")` can import excel files

# Reading in Data

Common ways to view data:

```
> head(sleepstudy)
```

	Reaction	Days	Subject
1	249.5600	0	308
2	258.7047	1	308
3	250.8006	2	308
4	321.4398	3	308
5	356.8519	4	308
6	414.6901	5	308



# Reading in Data

Common ways to view data:

```
> tail(sleepstudy)
```

	Reaction	Days	Subject
175	287.1726	4	372
176	329.6076	5	372
177	334.4818	6	372
178	343.2199	7	372
179	369.1417	8	372
180	364.1236	9	372

# Reading in Data

Common ways to view data:

```
> View(sleepstudy)  #In RStudio only
```

# Reading in Data

Common ways to view data:

```
> print(sleepstudy)
```

# Reading in Data

Common ways to view data:

```
> colnames(sleepstudy)
```

```
[1] "Reaction" "Days"      "Subject"
```

# Reading in Data

Common ways to view data:

```
> str(sleepstudy)
```

```
'data.frame':      180 obs. of  3 variables:
 $ Reaction: num  250 259 251 321 357 ...
 $ Days      : num  0 1 2 3 4 5 6 7 8 9 ...
 $ Subject   : Factor w/ 18 levels "308","309","310",...: 1 1
```

# Reading in Data

Common ways to view data:

```
> dim(sleepstudy)
```

```
[1] 180  3
```

# Reading in Data

Try it:

```
> blowdown = alr3::blowdown  
> ?blowdown
```

# Reading in Data

```
> summary(blowdown)
```

D	S	y	SP
Min. : 5.00	Min. :0.02175	Min. :0.0000	BS
1st Qu.: 9.00	1st Qu.:0.21856	1st Qu.:0.0000	BF
Median :14.00	Median :0.39326	Median :0.0000	JP
Mean :15.91	Mean :0.41156	Mean :0.4594	PB
3rd Qu.:21.00	3rd Qu.:0.58940	3rd Qu.:1.0000	A
Max. :85.00	Max. :0.98327	Max. :1.0000	C
			(Other)



# Reading in Data

```
> summary(blowdown$D)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
5.00	9.00	14.00	15.91	21.00	85.00

# Reading in Data

```
> with(blowdown,by(D,y,summary))
```

```
y: 0
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
5.00	7.00	11.00	13.09	16.00	85.00

---

```
y: 1
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
5.00	13.00	17.50	19.22	25.00	58.00

## Reading in Data

```
> colMeans(blowdown[, -4])
```

	D	S	y
	15.9050736	0.4115597	0.4593562

```
> #rowMeans()
```

# Reading in Data

```
> apply(blowdown[, -4], 2, mean)
```

D	S	y
15.9050736	0.4115597	0.4593562

# Reading in Data

```
> apply(blowdown[, -4], 2, function(x) mean(abs(x - mean(x))))
```

	D	S	y
	6.7392200	0.1965914	0.4966962

# Reading in Data

```
> blowdownScaled = scale(blowdown[, -4])
> apply(blowdownScaled, 2,
+       function(x) round(c(mean=mean(x), sd=sd(x)), 10))
```

```
      D S y
mean 0 0 0
sd    1 1 1
```

# Reading in Data

## Reshaping the data:

```
> require("reshape2")
> ###Long form:
> head(sleepstudy)
```

	Reaction	Days	Subject
1	249.5600	0	308
2	258.7047	1	308
3	250.8006	2	308
4	321.4398	3	308
5	356.8519	4	308
6	414.6901	5	308

# Reading in Data

## Reshaping the data:

```
> ###Wide form:
> sleepWide = acast(sleepstudy, Subject~factor(Days),
+                   value.var="Reaction")
> dim(sleepWide)
```

```
[1] 18 10
```

```
> head(round(sleepWide,1))
```

	0	1	2	3	4	5	6	7	8	9
308	249.6	258.7	250.8	321.4	356.9	414.7	382.2	290.1	430.6	466.4
309	222.7	205.3	203.0	204.7	207.7	216.0	213.6	217.7	224.3	237.3
310	199.1	194.3	234.3	232.8	229.3	220.5	235.4	255.8	261.0	247.5
330	321.5	300.4	283.9	285.1	285.8	297.6	280.2	318.3	305.3	354.0
331	287.6	285.0	301.8	320.1	316.3	293.3	290.1	334.8	293.7	371.6
332	234.9	242.8	273.0	309.8	317.5	310.0	454.2	346.8	330.3	253.9



# Reading in Data

## Reshaping the data:

```
> ###Back to tall form:
> sleepTall = melt(sleepWide,value.name="Reaction")
> dim(sleepTall)
```

```
[1] 180    3
```

```
> head(round(sleepTall,1))
```

	Var1	Var2	Reaction
1	308	0	249.6
2	309	0	222.7
3	310	0	199.1
4	330	0	321.5
5	331	0	287.6
6	332	0	234.9

## Writing data

```
> write.csv(sleepWide,file="sleepWide.csv")  
> write.table(sleepWide,file="sleepWide.txt",sep="\t")
```

# Linear Algebra

# Linear Algebra

Let's use the design matrix for sleepstudy:

```
> X = model.matrix(~Days, data=sleepstudy)
```

```
> class(X)
```

```
[1] "matrix"
```

```
> dim(X)
```

```
[1] 180  2
```

# Linear Algebra

Is the matrix of full rank?

```
> qr(X)$rank
```

```
[1] 2
```

# Linear Algebra

Matrix multiplication (and transposition)

```
> XtX = t(X)%*%X
```

```
> all.equal(XtX,crossprod(X,X))  #Uses less memory
```

```
[1] TRUE
```

# Linear Algebra

## Matrix Inversion

```
> XtXInv = solve(XtX)
> # XtXgenInv = MASS::ginv(Xtx)
> # XtXInv = chol2inv(chol(XtX))
```

# Linear Algebra

OLS estimates:

$$\hat{\beta} = (X'X)^{-1}X'Y$$

```
> XtXInv%%t(X)%%sleepstudy$Reaction
```

```
      [,1]
```

```
(Intercept) 251.40510
```

```
Days        10.46729
```

```
> lm(Reaction~Days,data=sleepstudy)$coef
```

```
(Intercept)      Days
```

```
251.40510      10.46729
```



# Linear Algebra

## US Judge ratings

A data frame containing 43 observations on 12 numeric variables.

```
[,1] CONT  Number of contacts of lawyer with judge.  
[,2] INTG  Judicial integrity.  
[,3] DMNR  Demeanor.  
[,4] DILG  Diligence.  
[,5] CFMG  Case flow managing.  
[,6] DECI  Prompt decisions.  
[,7] PREP  Preparation for trial.  
[,8] FAMI  Familiarity with law.  
[,9] ORAL  Sound oral rulings.  
[,10] WRIT  Sound written rulings.  
[,11] PHYS  Physical ability.  
[,12] RTEN  Worthy of retention.
```


```
> ?USJudgeRatings
```

# Linear Algebra

Principle Components (and (sloppy) biplot)

```
> judgesCent = scale(as.matrix(USJudgeRatings[,-1]),scale=F)
> Sigma = cov(judgesCent)
> eigs = eigen(Sigma)
> Scores = judgesCent %%% eigs$vectors[,1:2]
> plot(Scores,pch=16,ylim=c(-1,1)*2.5,xlab="",ylab="")
> arrows(0,0,eigs$vec[,1]*3,eigs$vec[,2]*3)
> text(eigs$vec[,1:2]*3,labels=colnames(judgesCent),
+      adj=c(1,0))
```

# Linear Algebra



# Linear Algebra

## Miscellaneous:


- kronecker product: `kronecker()` or `%x%`
- Hadamard Product: `A*B`
- Singular Value Decomposition: `svd()`
- Determinant or log determinant: `det()`, `determinant()`
- Diagonal, upper (or lower) triangle: `diag()`,  
`X[upper.tri(X)]`

# Data Visualization

# Univariate

## Data visualization: univariate

```
> yMeans = as.numeric(with(blowdown,by(y,SPP,mean)))  
> barplot(yMeans,names.arg=levels(blowdown$SPP))
```





## Data visualization: univariate

```
> par(mfrow=c(1,2))  
> pie(yMeans,labels=levels(blowdown$SPP),  
+     main="Survival rates")  
> plotrix::pie3D(yMeans,labels=levels(blowdown$SPP),  
+               explode=0.1,  
+               col=rgb(20/256,c(1:9*20)/256,120/256))  
> par(mfrow=c(1,1))
```




# Data visualization: univariate




## Data visualization: univariate

```
> boxplot(blowdown$D)
```



## Data visualization: univariate


```
> boxplot(blowdown$D,main="Blowdown data",ylab="Diameter")
```




## Data visualization: univariate

```
> par(mfrow=c(1,2))  
> with(blowdown,by(D,y,boxplot,main="Blowdown data",  
+                               ylab="Diameter"))  
> par(mfrow=c(1,1))
```

**Blowdown data**




**Blowdown data**



## Data visualization: univariate

```
> boxplot(blowdown$D~blowdown$y,main="Blowdown data",  
+         ylab="Diameter",names=c("survive","died"),  
+         cex.lab=1.5,cex.axis=1.5)
```


**Blowdown data**



## Data visualization: univariate

```
> hist(blowdown$D,main="Blowdown data",xlab="",  
+       ylab="Diameter",cex.lab=1.5,cex.axis=1.5)
```

**Blowdown data**




## Data visualization: univariate

```
> plot(density(blowdown$D),ylab="Diameter",cex.lab=1.5,  
+      cex.axis=1.5,main="blowdown data")
```

## Data visualization: univariate

**blowdown data**



N = 3666 Bandwidth = 1.496




## Data visualization: univariate

```
> hist(blowdown$D,main="Blowdown data",xlab="",freq=FALSE,  
+       ylab="Diameter",cex.lab=1.5,cex.axis=1.5)  
> lines(density(blowdown$D,bw=2.5),col="red",lwd=2)
```

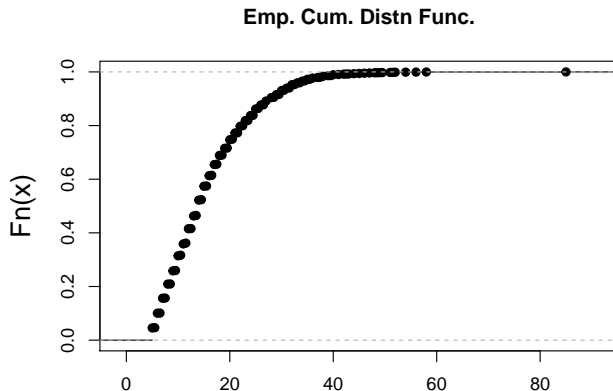
# Data visualization: univariate

**Blowdown data**



## Data visualization: univariate

```
> plot(ecdf(blowdown$D),main="Emp. Cum. Distn Func.",  
+       xlab="",cex.lab=1.5)
```



## Data visualization: univariate

**Plotting a custom function:** Example: Dose-response  
Exponential model


$$\text{Prob}(\text{Infection}) = 1 - e^{-rd}$$

```
> DRfun = function(dose){  
+   return( 1-exp(-2.18E-04*dose) )  
+ }
```

## Data visualization: univariate

**Plotting a custom function:** Example: Dose-response  
Exponential model

$$Prob(Infection) = 1 - e^{-rd}$$




# Data visualization: univariate

## qq-plots


```
> set.seed(1)
> x1 = rnorm(100)
> x2 = rgamma(100,5,10)
> par(mfrow=c(1,2))
> qqnorm(x1);qqline(x1)
> qqnorm(x2);qqline(x2)
```

# Data visualization: univariate

Normal Q-Q Plot



Normal Q-Q Plot





# Data visualization: univariate

## qq-plots

```
> par(mfrow=c(1,2))  
> qqplot(qgamma(ppoints(500),5,10),x1)  
> qqline(x1,distribution=function(p)qgamma(p,5,10))  
> qqplot(qgamma(ppoints(500),5,10),x2)  
> qqline(x2,distribution=function(p)qgamma(p,5,10))
```



# Data visualization: univariate




# Data visualization: univariate

## Time series:


Example: Average monthly deaths from lung diseases in the UK from 1974-1979

```
> class(ldeaths)
> plot(ldeaths,ylab="# Deaths")
> plot(stl(ldeaths,s.window="periodic"))
```

# Data visualization: univariate



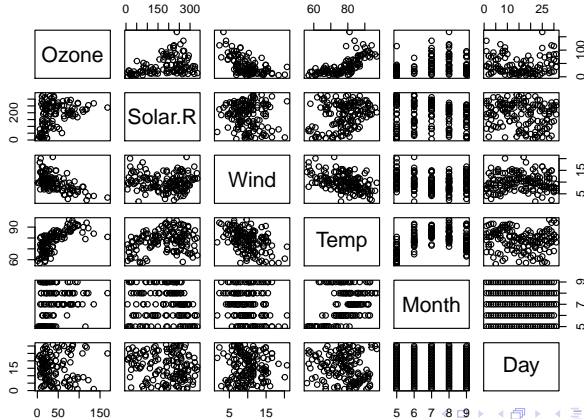
# Data visualization: univariate



# Multivariate

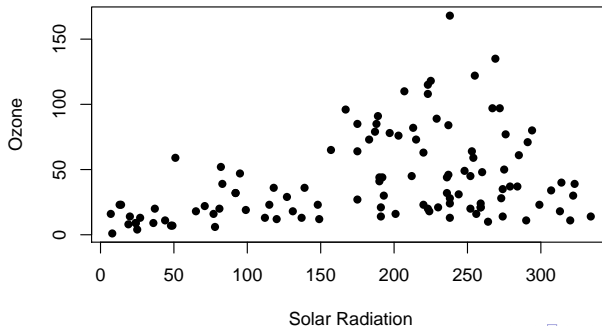
# Data visualization: multivariate

```
> pairs(airquality)
```



## Data visualization: multivariate

```
> plot(Ozone~Solar.R,data=airquality,xlab="Solar Radiation"  
+      ylab="Ozone",pch=16)
```

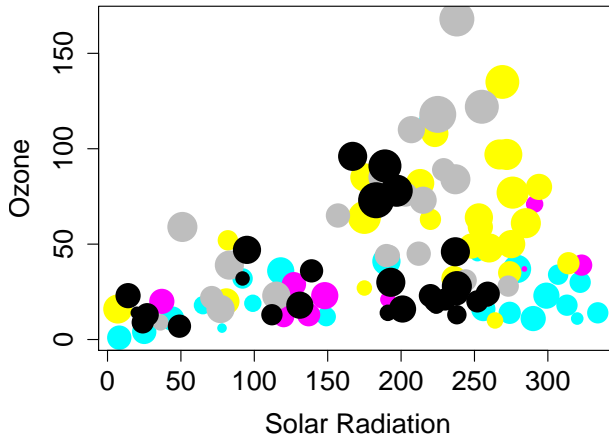


# Data visualization: multivariate

```
> CEXs = seq(0.5,5,length.out=500)
> CEXs = with(airquality,
+             CEXs[length(CEXs)*(max(Wind)+1-Wind)/
+             (max(Wind)+1)])
> with(airquality,
+      plot(Ozone~Solar.R,xlab="Solar Radiation",
+      ylab="Ozone",pch=16,cex=CEXs,col=Month,
+      cex.lab=1.5,cex.axis=1.5))
```




## Data visualization: multivariate



## Data visualization: multivariate

```
> with(airquality, sapply(unique(Month),
+   function(x){ind=which(Month==x);
+     abline(lm(Ozone[ind]~Solar.R[ind]),
+       col=x, lwd=3)}))
> with(airquality, legend("topleft", lwd=rep(4,5), cex=1.5,
+   col=unique(Month), legend=unique(Month)))
```


## Data visualization: multivariate



# Data visualization: multivariate

```
> plot(sleepTall$Reaction~sleepTall$Var2,pch=16,
+       xlab="Days of sleep deprivation",
+       ylab="Reaction time",cex.lab=1.5,cex.axis=1.5)
> for(i in 1:nrow(sleepWide)){
+   lines(sleepWide[i,]~c(0:9))
+ }
> lines(colMeans(sleepWide)~c(0:9),col="blue",lwd=4,lty=2)
```


## Data visualization: multivariate



# Data visualization

```
> curve(dnorm(x,-1.5),-4.5,3.5)
> xseq = seq(0,3.5,length.out=500)
> yseq = dnorm(xseq,-1.5)
> polygon(x=c(xseq,xseq[length(xseq):1]),0),
+         y=c(yseq,rep(0,length(xseq)+1)),
+         col=rgb(0.25,0.75,1,alpha=0.5))
```

# Data visualization



# Data visualization

Other miscellaneous functions:

```
> text()  
> segments()  
> arrows()  
> symbols()  
etc. . .
```



# Data visualization

You can save to disk your images:

```
> jpeg()  
> pdf()  
> png()  
> bmp()  
> tiff()
```

# Data visualization

You can save to disk your images:

```
> jpeg()  
> pdf()  
> png()  
> bmp()  
> tiff()
```

Syntax is something like:

```
> jpeg("foo.jpg",height=800,width=800)#in pixels  
> plot( ... )  
> dev.off()
```

# ggplot2


# Data visualization: ggplot2

**ggplot2:** ( <http://docs.ggplot2.org/current/> )

Syntax looks considerably different!

```
> require("ggplot2")
> blowdown$y <-
+   factor(blowdown$y, labels=c("Surv", "Died"))
> ggplot(blowdown, aes(factor(y), D)) +
+   geom_boxplot() +
+   theme(axis.text=element_text(size=20),
+         axis.title=element_text(size=0, colour="white"))
```


## Data visualization: ggplot2



# Data visualization: ggplot2

```
> ggplot(blowdown, aes(D, fill=SPP)) +  
+   geom_histogram() +  
+   theme(axis.text=element_text(size=20),  
+         axis.title=element_text(size=0, colour="white"))
```

# Data visualization: ggplot2




## Data visualization: ggplot2

```
> ggplot(blowdown, aes(D, fill=SPP, colour=SPP)) +  
+   geom_density(alpha=0.1) +  
+   theme(axis.text=element_text(size=20),  
+         axis.title=element_text(size=0, colour="white"))
```




# Data visualization: ggplot2



## Data visualization: ggplot2

```
> ggplot(blowdown, aes(D, fill=SPP, colour=SPP)) +  
+   geom_density(position="stack") +  
+   theme(axis.text=element_text(size=20),  
+         axis.title=element_text(size=0, colour="white"))
```


# Data visualization: ggplot2



## Data visualization: ggplot2

```
> ggplot(airquality, aes(Solar.R, Ozone)) +  
+   geom_point(size=6, aes(color=Temp)) +  
+   geom_smooth(method="lm") +  
+   theme(axis.text=element_text(size=20),  
+         axis.title=element_text(size=0, colour="white"))
```

# Data visualization: ggplot2



# 3D plotting

## Data visualization: 3D


- `scatterplot3d`
- `scatter3d`
- `plot3D` / `plot3Drgl`
- `emdbook`

# Data visualization: 3D


```
> normMix = function(x,y){
+   mixProbs = c(1/3,1/2,1/6)
+   ret = dnorm(x,-2)*dnorm(y,-2) +
+     dnorm(x,0)*dnorm(y,2) +
+     dnorm(x,2)*dnorm(y,0)
+   return(ret)
+ }
> curve3d(normMix(x,y),from=c(-5,-5),to=c(5,5),sys3d="contour",
+   xlab="",ylab="",labcex=1.5,nlevels=20)
> curve3d(normMix(x,y),from=c(-5,-5),to=c(5,5),sys3d="persp",theta=-15,
+   xlab="",ylab="",zlab="")
> curve3d(normMix(x,y),from=c(-5,-5),to=c(5,5),sys3d="rgl",
+   xlab="",ylab="",zlab="",
+   col = rgb(20/256,60/256,120/256,0.5))
```




## Data visualization: 3D



# Data visualization: 3D



## Data visualization: 3D




## Data visualization: 3D

Hydrocarbon vapor pollution (g) vs. tank temperature (F) and and tank pressure (psi)

```
> require("plot3D");require("alr3")
> pollution = sniffer[,c("TankTemp","TankPres")]
> pollution = data.frame(pollution,logy=log(sniffer$Y))
> with(pollution,scatter3D(TankTemp,TankPres,logy,theta=45,
+                           phi=20,xlab="Temp",ylab="Press",
+                           zlab="Pollution",pch=16))
```


## Data visualization: 3D



## Data visualization: 3D

```
> with(pollution, scatter3D(TankTemp, TankPres, logy, theta=45,  
+                             phi=20, xlab="Temp", ylab="Press",  
+                             zlab="Pollution", pch=16,  
+                             type="h"))
```

# Data visualization: 3D




## Data visualization: 3D

```
> fit = lm(logy~TankTemp+TankPres,data=pollution)
> xgrid=with(pollution,
+           seq(min(TankTemp),max(TankTemp),length.out=15))
> ygrid=with(pollution,
+           seq(min(TankPres),max(TankPres),length.out=15))
> xygrid = expand.grid(TankTemp=xgrid,TankPres=ygrid)
> logyPred = matrix(predict(fit,xygrid),15,15)
> fitPts = predict(fit)
> with(pollution,
+     scatter3D(TankTemp,TankPres,logy,theta=45,
+               phi=20,xlab="Temp",ylab="Press",
+               zlab="Pollution",pch=16,
+               surf=list(x=xgrid,y=ygrid,z=logyPred,
+               fit=fitPts,facets=NA)))
```




## Data visualization: 3D




## Data visualization: 3D

```
> BDTab = table(cut(blowdown$D,seq(5,85,by=10),  
+               include.lowest=TRUE),  
+               blowdown$SPP)  
> hist3D(z=BDTab,col = rgb(20/256,60/256,120/256),  
+        border = "black",shade = 0.4,space = 0.15,  
+        xlab="Diameter",ylab="Species",zlab="")  
> plotrgl()
```

# Data visualization: 3D



## Data visualization: 3D



# Data visualization

Lots of other fun toys:

- Spatial visualization: Rgooglemaps, rworldmap, tmap, etc.
- Network visualization: igraph, networkDynamic, etc.
- animation
- Many others!!!

# (Legal) Performance Enhancers

# Rcpp

We'll focus on compiler, foreach, and Rcpp (with inline).  
 Our (trivial) example function will be

$$f(n) = \sum_{i=1}^n i \quad \left( = \frac{n(n+1)}{2} \right)$$

```
> RFun = function(n,start=1){
+   ret=0
+   for(i in start:n){
+     ret = ret + i
+   }
+   return(ret)
+ }
> RFun(100) == 100*101/2

[1] TRUE
```

# compiler

Thanks to our own Luke Tierney!

```
> require("compiler")
> RFunComp = cmpfun(RFun)
> system.time(RFun(1e7))

  user  system elapsed 
 3.47    0.00    3.48 

> system.time(RFunComp(1e7))

  user  system elapsed 
 0.35    0.00    0.35
```



## foreach (parallelizing)

If we have an embarassingly parallel task, use the foreach function.

```
> require("foreach")
> require("doParallel")
> registerDoParallel(cl=2)
> system.time({
+   temp = foreach(i=c(5e6,1e7)) %dopar% {
+     if(i==5e6){
+       RFun(i)
+     }else{
+       RFun(i,5e6+1)
+     }
+   }
+   print(sum(unlist(temp)) == 1e7*(1e7+1)/2)
+ })
```

```
[1] TRUE
```

```
user system elapsed
0.02    0.00    2.81
```

```
> stopImplicitCluster()
```

# Rcpp

Now in c++:

```
> require("Rcpp");require("inline")
> cppFun = cxxfunction(signature(nFromR="integer"),
+                        body='
+ int n = Rcpp::as<int>(nFromR);
+ int ret = 0;
+ for(int i=1;i<n+1;i++){
+   ret += i;
+ }
+ return wrap(ret);
+ ',plugin="Rcpp")
> cppFun(100)

[1] 5050
```

# Rcpp

```
> system.time(RFun(1e7))  
  
   user   system elapsed  
  3.58    0.02    3.62  
  
> system.time(cppFun(1e7))  
  
   user   system elapsed  
  0.02    0.00    0.01
```

# Rcpp

Lots of additional Rcpp packages, such as RcppSugar,  
RcppArmadillo, etc.

# Extended Examples

# Blowdown data

## Extended example on the blowdown data

Data from the Boundary Waters Canoe Area Wilderness Blowdown.  
 The data frame blowdown includes nine species of trees.

```
> attach(blowdown)
```

```
> dim(blowdown)
```

```
[1] 3666    4
```

```
> head(blowdown)
```

	D	S	y	SPP
1	9	0.0217509	Surv	BA
2	14	0.0217509	Surv	BA
3	18	0.0217509	Surv	BA
4	23	0.0217509	Surv	BA
5	9	0.0217509	Surv	BA
6	16	0.0217509	Surv	BA

## Extended example on the blowdown data

Does species influence survival?

```
> y = factor(y, labels=c("Survived", "Died"))
> ( tab = table(SPP, y) )
```

	y	
SPP	Survived	Died
A	130	306
BA	69	6
BF	426	233
BS	438	532
C	311	44
JP	89	413
PB	407	90
RM	101	22
RP	11	38



## Extended example on the blowdown data

Does species influence survival?

```
> chisq.test(tab)
```

Pearson's Chi-squared test

```
data: tab
```

```
X-squared = 848.72, df = 8, p-value < 2.2e-16
```

```
> prop.table(tab,1)
```

y

SPP	Survived	Died
A	0.2981651	0.7018349
BA	0.9200000	0.0800000
BF	0.6464340	0.3535660
BS	0.4515464	0.5484536
C	0.8760563	0.1239437

## Extended example on the blowdown data

```
> image(t(prop.table(tab,1)),xaxt="n",yaxt="n",
+       main="Tree Species and Survival")
> box()
> axis(1,at=0:1,labels=levels(y),cex.axis=1.5)
> axis(2,at=seq(1,0,length.out=length(levels(SPP))),
+       labels=levels(SPP),las=1,cex.axis=1.5)
> text(cbind(rep(0:1,each=length(levels(SPP))),
+            rep(seq(0,1,length.out=length(levels(SPP))),2)),
+       labels=round(c(prop.table(tab,1)),2),cex=1.5)
```

## Extended example on the blowdown data

Tree Species and Survival

A	0.22	0.78
BA	0.82	0.18
BF	0.82	0.18
BS	0.18	0.82
C	0.88	0.12
JP	0.45	0.55
PB	0.65	0.35
RM	0.92	0.08
RP	0.3	0.7
	Survived	Died

## Extended example on the blowdown data

How about the diameter?

```
> logD = log(D)
> t.test(logD~y,var.equal=FALSE)
```

Welch Two Sample t-test

data: logD by y

t = -27.062, df = 3663.9, p-value < 2.2e-16

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

-0.4645252 -0.4017640

sample estimates:

mean in group Survived  
2.427742



mean in group Died  
2.860887

## Extended example on the blowdown data

Are the results valid? I.e., are they normal? Probably not. . .

```
> par(mfrow=c(1,2))  
> by(logD,y,qqnorm)  
> par(mfrow=c(1,1))
```

# Extended example on the blowdown data



## Extended example on the blowdown data

Let's do a non-parametric test to verify the association:

```
> kruskal.test(logD,y)
```

Kruskal-Wallis rank sum test

data: logD and y

Kruskal-Wallis chi-squared = 598.53, df = 1, p-value < 2.2e-16


## Extended example on the blowdown data

Density plots by survival:

```
> ind=which(y=="Died")  
> plot(density(x=logD[ind]),ylim=c(0,0.9),  
+      main="log(Diameter)")  
> lines(density(x=logD[-ind]),lty=2)  
> legend("topright",legend=c("Died","Survived"),lty=1:2)
```



## Extended example on the blowdown data




N = 1684 Bandwidth = 0.09041

## Extended example on the blowdown data

So there likely exists this confounding factor in that species have differing avg. diameters.

```
> boxplot(logD~SPP)
```



## Extended example on the blowdown data

Is species related to diameter?

```
> lmMod= lm(logD~SPP)
> anova(lmMod)
```

Analysis of Variance Table

Response: logD

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
SPP	8	405.36	50.670	287.95	< 2.2e-16 ***
Residuals	3657	643.52	0.176		

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

# Extended example on the blowdown data

```
> summary(lmMod)

Call:
lm(formula = logD ~ SPP)

Residuals:
    Min       1Q   Median       3Q      Max
-1.47181 -0.26516  0.01793  0.29341  1.36408

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   3.08125     0.02009  153.374 < 2e-16 ***
SPPBA         -0.77845     0.05244  -14.845 < 2e-16 ***
SPPBF         -0.90196     0.02590  -34.829 < 2e-16 ***
SPPBS         -0.51351     0.02419  -21.231 < 2e-16 ***
SPPC          -0.43022     0.02999  -14.346 < 2e-16 ***
SPPJP          0.01014     0.02746   0.369  0.712
SPPPB         -0.63780     0.02753  -23.171 < 2e-16 ***
SPPRM         -0.60928     0.04283  -14.226 < 2e-16 ***
SPPRP          0.49170     0.06320   7.780 9.4e-15 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4195 on 3657 degrees of freedom
Multiple R-squared:  0.3865,    Adjusted R-squared:  0.3851
F-statistic: 287.9 on 8 and 3657 DF, p-value: < 2.2e-16
```


## Extended example on the blowdown data

Let's check the model assumptions:


```
> qqnorm(resid(lmMod));qqline(resid(lmMod))  
> plot(resid(lmMod)~fitted(lmMod));abline(h=0)
```

## Extended example on the blowdown data

Normal Q-Q Plot



## Extended example on the blowdown data




## Extended example on the blowdown data

Another factor to consider is the severity of the storm

```
> plot(density(x=S),ylim=c(0,2.25),main="Storm Severity")  
> lines(density(x=S[-ind]),lty=2)  
> legend("topright",legend=c("Died","Survived"),lty=1:2)
```



## Extended example on the blowdown data



N = 3666 Bandwidth = 0.04038

## Extended example on the blowdown data

A more appropriate model would be logistic regression

```
> ( logRegMod = glm(y~logD + SPP + S, family=binomial) )
```

```
Call: glm(formula = y ~ logD + SPP + S, family = binomial)
```

Coefficients:

(Intercept)	logD	SPPBA	SPPBF	SPPBS
-5.9971951	1.5813423	-2.2427869	0.0002284	0.1672262
SPPJP	SPPPB	SPPRM	SPPRP	S
1.0399651	-1.7235679	-1.7956738	0.0031381	4.6288861

Degrees of Freedom: 3665 Total (i.e. Null); 3655 Residual

Null Deviance: 5058

Residual Deviance: 3259 AIC: 3281

## Extended example on the blowdown data

```
> round(summary(logRegMod)$coef,4)
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-5.9972	0.3748	-15.9993	0.0000
logD	1.5813	0.1115	14.1876	0.0000
SPPBA	-2.2428	0.4936	-4.5439	0.0000
SPPBF	0.0002	0.1789	0.0013	0.9990
SPPBS	0.1672	0.1518	1.1020	0.2705
SPPC	-2.0765	0.2162	-9.6031	0.0000
SPPJP	1.0400	0.1788	5.8176	0.0000
SPPPB	-1.7236	0.1865	-9.2435	0.0000
SPPRM	-1.7957	0.3019	-5.9472	0.0000
SPPRP	0.0031	0.4132	0.0076	0.9939
S	4.6289	0.2128	21.7477	0.0000

## Extended example on the blowdown data

We could also have done a probit regression model

```
> ( probRegMod = glm(y~logD + SPP + S,
+                    family=binomial(link="probit"))) )
```

```
Call: glm(formula = y ~ logD + SPP + S, family = binomial(link
```

Coefficients:

(Intercept)	logD	SPPBA	SPPBF	SPPBS
-3.350235	0.875785	-1.186295	-0.031899	0.091554
SPPJP	SPPPB	SPPRM	SPPRP	S
0.600456	-0.971438	-0.966891	-0.005013	2.652496

Degrees of Freedom: 3665 Total (i.e. Null); 3655 Residual

Null Deviance: 5058

Residual Deviance: 3281 AIC: 3303

## Extended example on the blowdown data

We could also have done a probit regression model

```
> summary(probRegMod)$coef
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-3.35023538	0.20911836	-16.02076157	9.152243e-58
logD	0.87578494	0.06240703	14.03343517	9.732093e-45
SPPBA	-1.18629531	0.26124554	-4.54092074	5.600909e-06
SPPBF	-0.03189922	0.10327689	-0.30887084	7.574198e-01
SPPBS	0.09155382	0.08782421	1.04246677	2.971953e-01
SPPC	-1.17118254	0.11825869	-9.90356469	4.016963e-23
SPPJP	0.60045604	0.10104811	5.94227897	2.810865e-09
SPPPB	-0.97143796	0.10571710	-9.18903323	3.963877e-20
SPPRM	-0.96689104	0.16695204	-5.79142987	6.978970e-09
SPPRP	-0.00501280	0.23671805	-0.02117625	9.831051e-01
S	2.65249622	0.11714384	22.64307109	1.632635e-113


## Extended example on the blowdown data

We can also plot the probability surface for a particular species, say, black spruce.

```
> logitFun = function(x,y){
+   eta = logRegMod$coef["(Intercept)"] +
+     logRegMod$coef["SPPBS"]+logRegMod$coef["logD"]*x+
+     logRegMod$coef["S"]*y
+   return(1/(1+exp(-eta)))
+ }
> with(blowdown,
+   curve3d(logitFun,from=c(min(logD),min(S)),
+     to=c(max(logD),max(S)),sys3d="persp",
+     theta=-20,xlab="Diameter",ylab="Storm",
+     zlab="Probability of Dying"))
```

## Extended example on the blowdown data

We can also plot the probability surface for a particular species, say, black spruce.



# UK monthly deaths from lung diseases




## Extended example on the lung diseases data

Consider monthly deaths from bronchitis, emphysema and asthma in the UK, 1974-1979.

```
> yr <- floor(tt <- time(mdeaths))
> plot(ldeaths,ylab="",
+      xy.labels=paste(month.abb[12*(tt-yr)],
+                      yr-1900,sep=" ' "))
```

## Extended example on the lung diseases data



## Extended example on the lung diseases data

Specifically we want to get predicted values for future numbers of lung disease fatalities.

The statistical model is

$$y_t = \phi y_{t-1} + \epsilon_t$$

where  $\epsilon_t \stackrel{iid}{\sim} N(0, \sigma^2)$ . We'll construct a Gibbs sampler (a MCMC algorithm). This can be done by iteratively drawing from

$$\pi(\sigma^2 | data)$$

$$\pi(\psi | \sigma^2, data)$$

$$\pi(y_{new} | \psi, \sigma^2, data)$$

## Extended example on the lung diseases data

These conditional distributions (with uninformative improper priors) are given as follows. First, let

$$B = \sum_t y_{t-1}^2$$

$$b = \sum_t y_t y_{t-1} / B$$

$$Q_b = \sum_t (y_t - b y_{t-1})^2.$$

Then we have

$$\pi(\sigma^2 | \phi, \text{data}) \sim \Gamma^{-1}(0.5(T - 2), 0.5Q_b)$$

$$\pi(\phi | \sigma^2, \text{data}) \sim N(b, \sigma^2 / B)$$

$$\pi(y_{\text{new}} | \phi, \sigma^2, \text{data}) \sim N(\phi y_T, \sigma^2).$$

## Extended example on the lung diseases data

Now to do this in R.

```
> require("MCMCpack")
> fun1= function(y,nsims=1000,stepsAhead=100){
+   TT = length(y)
+   BB = sum(y[-TT]^2)
+   bb = sum(y[-1]*y[-TT])/BB
+   Qb = sum((y[-1]-bb*y[-TT])^2)
+
+   s2 = rinvgamma(nsims,shape=0.5*(TT-3),scale=0.5*Qb)
+   phi = rnorm(nsims,mean=bb,sd=sqrt(s2/BB))
+   ypred = matrix(0.0,nsims,stepsAhead)
+   ypred[,1]=rnorm(nsims,phi*y[TT],sd=sqrt(s2))
+   for(tt in 2:stepsAhead){
+     ypred[,tt] = rnorm(nsims,phi*ypred[,tt-1],sd=sqrt(s2))
+   }
+   return(list(s2=s2,phi=phi,ypred=ypred))
+ }
```

## Extended example on the lung diseases data


Let's test the Bayesian estimators out via simulation:

```
> set.seed(1)
> M=1000
> phiTrue = 0.75
> s2True = 0.5
> TT=500
> Y = matrix(NA,M,TT+1)
> Y[,1]=rnorm(M,sd=sqrt(s2True/(1-phiTrue^2)))
> phiHat = s2Hat = numeric(M)
> for(tt in 1+1:TT){
+   Y[,tt] = 0.75*Y[,tt-1] + rnorm(M,sd=sqrt(s2True))
+ }
> for(iter in 1:M){
+   temp = fun1(Y[iter,],stepsAhead=2)
+   s2Hat[iter] = mean(temp$s2)
+   phiHat[iter] = mean(temp$phi)
+ }
```

## Extended example on the lung diseases data


```
> hist(phiHat)
> abline(v=phiTrue,col="red",lwd=2)
> hist(s2Hat)
> abline(v=s2True,col="red",lwd=2)
```

## Extended example on the lung diseases data





# Extended example on the lung diseases data




## Extended example on the lung diseases data

Now let's apply this to our lung disease mortality data:

```
> ldeaths1 = ldeaths - mean(ldeaths)
> fit1 = fun1(ldeaths1, nsims=1e5, stepsAhead=12)
> ypreds = colMeans(fit1$ypred) + mean(ldeaths)
> yBounds = t(apply(fit1$ypred, 2, quantile,
+                   probs=c(0.025, 0.975))) + mean(ldeaths)
> YL = range(c(ldeaths, yBounds))
> plot(c(ldeaths), ylab="", xlab="", type="l",
+       xlim=c(1, (length(ldeaths)+12)), ylim=YL, lwd=2)
> lines(ypreds ~ c(length(ldeaths)+1:12), col="red", lwd=2)
> lines(yBounds[,1] ~ c(length(ldeaths)+1:12), col="blue",
+       lwd=2, lty=2)
> lines(yBounds[,2] ~ c(length(ldeaths)+1:12), col="blue",
+       lwd=2, lty=2)
```

## Extended example on the lung diseases data


Now let's apply this to our lung disease mortality data:




## Extended example on the lung diseases data

```
> TT = length(ldeaths1); BB = sum(ldeaths1[-TT]^2);
> bb = sum(ldeaths1[-1]*ldeaths1[-TT])/BB;
> Qb = sum((ldeaths1[-1]-bb*ldeaths1[-TT])^2)
> curve(dgamma(x,shape=0.5*(TT-3),scale=0.5*Qb),
+       lwd=2,xlab=expression(sigma^2),from=0,to=5e8,
+       ylab="",cex.axis=1.5,cex.lab=1.5)
> curve(dnorm(x,bb,sd=sqrt(mean(fit1$s2)/BB)),
+       lwd=2,xlab=expression(phi),from=0,to=1,
+       ylab="",cex.axis=1.5,cex.lab=1.5)
```

## Extended example on the lung diseases data



## Extended example on the lung diseases data



## Extended example on the lung diseases data

Can we make it faster?

```
> require("compiler")
> fun1Comp = cmpfun(fun1)
```

```
funcpp = cxxfunction(
  signature(YY="numeric",NSIMS="integer",STEPSAHEAD="integer"),
  body='
```

```
//Environment stats("package:stats");
//Function rnorm = stats["rnorm"];
Environment MCMCpack("package:MCMCpack");
Function rinvgamma = MCMCpack["rinvgamma"];

...
```

## Extended example on the lung diseases data

Let's compare:

```
> system.time({for(it in 1:100){
+   Rsims = fun1(Y[1,])
+   }
+ })
```

```
user system elapsed
1.34    0.00    1.35
```

```
> system.time({for(it in 1:100){
+   RsimsComp = fun1Comp(Y[1,])
+   }
+ })
```

```
user system elapsed
1.56    0.00    1.59
```

```
> system.time({for(it in 1:100){
+   Rcppsims = funcpp(Y, NSIMS=1000, STEPSAHEAD=100)
+   }
+ })
```

```
user system elapsed
0.89    0.08    0.97
```



# THE END

